ANMOL DEEP

AI Engineer

Bengaluru, India · adsk2050@gmail.com · linkedin.com/in/anmoldeep1 · github.com/adsk2050 · adsk2050.github.io

Summary

AI engineer, 6 years, IIT Guwahati. Three years building LLM systems end to end: multi-agent systems, RAG pipelines, voice agents, and the evaluation and guardrails around them. Earlier, trained, deployed and monitored machine learning models on documents, images and customer data.

Experience

Independent projects · *sole engineer* *Bengaluru* · Sep 2025 – present

- Shrota, an AI phone-survey platform. A researcher designs a survey in a web builder, and a voice agent calls each respondent and asks the questions in Hindi. Built all of it alone. Pre-launch, run on test calls.
- Measured the voice agent's response delay stage by stage and is cutting it. Latency and cost are tracked on every call, and calls carry on with a backup speech or LLM vendor when one fails.
- Research, not published: had an LLM rewrite only the changed part of a schema-governed output. In experiments on Gemini, edits took 172x fewer output tokens than regenerating the whole output (github.com/adsk2050/research-experiments).

Machine Learning Engineer · *Coreworks AI* *Bengaluru* · Mar 2026 – Aug 2026

- Owned the correctness of agent output on a multi-agent product that turns business data into traceable reports, decks and spreadsheet models. Built the agent that writes the spreadsheets as structured output.
- Set up an evaluation pipeline for agent output with Langfuse, and added input and security guardrails.
- Built, almost alone, a server that lets LLM agents author slide decks over MCP and REST APIs.

Product Engineer · *Aidora Inc (via Multiplier)* *Remote* · Aug 2024 – Sep 2025

- Built voice agents on LiveKit that callers can interrupt mid-sentence, and a multi-agent system that reaches internal APIs through typed tools over MCP.
- Built a retrieval-augmented generation (RAG) pipeline: a crawler and document ingestion that feed Milvus vector search, with query rewriting and reranking.
- Built the system that evaluates the team's text and voice agents, and wrote the system design documents. Trained two incoming engineers.

NLP Engineer · *Staram Analytics* *Remote* · May 2024 – Aug 2024

- Built a knowledge graph of Indian law in Neo4j from thousands of scanned legal documents, with retrieval over the graph, and an agent that computes accident compensation from court guidelines.

NLP Engineer · *Insurance Samadhan* *Delhi NCR* · Sep 2023 – May 2024

- Trained, deployed and monitored models that classify insurance documents and extract key fields from them. Built the annotation pipeline for the training data, and the model evaluation.
- Readied the system for an order of 60,000+ documents with a cost analysis of each part, and presented compute options to the CTO.

Senior Associate · *I-PAC* *Bihar* · Sep 2022 – Aug 2023

- Built Hindi text extraction from scanned public records with OpenCV and Tesseract, and turned data analysis into outreach plans.

Data Scientist · *Tata Digital (Tata CLiQ)* *Mumbai / Remote* · Aug 2020 – Sep 2022

- Built product recommenders from image embeddings and nearest-neighbour search, and churn prediction. Promoted after one year, then led two graduate hires.

Education

IIT Guwahati · B.Tech Mechanical Engineering, Minor in Product Design Jul 2016 – Jul 2020

Skills

LLM systems: RAG, multi-agent systems, tool calling, structured outputs, LangGraph, LangChain, MCP, evaluation and observability (Langfuse, Ragas, OpenTelemetry), guardrails
Voice and ML: LiveKit, speech-to-text, text-to-speech, turn detection, Milvus, FAISS, Neo4j, PyTorch, Hugging Face, scikit-learn, spaCy, OpenCV
Backend and cloud: Python, FastAPI, PostgreSQL, Redis, Docker, Kubernetes, AWS, GCP, GitHub Actions